

# SE 2228-Analysis and Design of Algorithms

## Algorithm Analysis-2

### Iterative/Recursive Algorithms

# Computing and Simplifying Running Time

- Method: Associate a "cost" with each statement and find the "total cost" by finding the total number of times each statement is executed.
- Express running (Execution) time in terms of the size of the problem.
- Although some algorithms may have different cost values, their asymptotic complexity expression may be the same.

Example  $T1(n) = c1*n$  and  $T2(n) = (c1+c2)*n$

Both are  $O(n)$  algorithms.

# General Rules for Running Time Estimation

- **Loops:** The running time of a loop is at most the running time of the **statements inside of that loop times the number of iterations.**
- **Nested Loops:** Running time of a nested loop containing a statement in the inner most loop is the running time of statement multiplied by the product of the sizes of all loops.
- **Consecutive Statements:** Just add the running times of those consecutive statements.
- **If/Else:** Never more than the running time of the test plus the larger of running times of S1 and S2.

# Computing Execution Time of Algorithms

- Each operation in an algorithm (or a program) has a cost.  
→ Each operation takes a certain of time.

count = count + 1;

takes a certain amount of time, but it is constant

A sequence of operations:

count = count + 1;                      Cost:  $c_1$

sum = sum + count;                      Cost:  $c_2$

Total Cost =  $c_1 + c_2$

# Computing Execution Time of Algorithms

Example: Simple If-Statement

	<u>Cost</u>	<u>Times</u>
if (n < 0)	c1	1
absval = -n	c2	1
else		
absval = n;	c3	1

**Total Cost**  $\leq c1 + \max(c2, c3)$

## Example1: Simple Loop

	<u>Cost</u>	<u>Times</u>
<code>i = 1;</code>	c1	1
<code>sum = 0;</code>	c2	1
<code>while (i &lt;= n) {</code>	c3	n+1
<code>i = i + 1;</code>	c4	n
<code>sum = sum + i;</code>	c5	n
<code>}</code>		

$$\text{Total Cost} = c1 + c2 + (n+1)*c3 + n*c4 + n*c5$$

➔ The time required for this algorithm is proportional to n

# Computing Execution Time of Algorithms

## Example-2:Nested loop

	<u>Cost</u>	<u>Times</u>
i=1;	c1	1
sum = 0;	c2	1
while (i <= n) {	c3	n+1
j=1;	c4	n
while (j <= n) {	c5	n*(n+1)
sum = sum + i;	c6	n*n
j = j + 1;	c7	n*n
}		
i = i + 1;	c8	n
}		

$$\begin{aligned}T(n) &= c1 + c2 + (n+1)*c3 + n*c4 + n*(n+1)*c5 + n*n*c6 + n*n*c7 + n*c8 \\&= (c5+c6+c7)*n^2 + (c3+c4+c5+c8)*n + (c1+c2+c3) \\&= a*n^2 + b*n + c\end{aligned}$$

➔ The growth-rate function for this algorithm is **O(n<sup>2</sup>)**

# Asymptotic Analysis

- Using *rate of growth* as a measure to compare different functions implies comparing them asymptotically.
- If  $f(x)$  is *growing faster* than  $g(x)$ , then  $f(x)$  always eventually becomes larger than  $g(x)$  *in the limit* (for large enough values of  $x$ ).

How to express rate of growth for different algorithm types?

We consider two main categories of algorithms:

- *Iterative (nonrecursive)* Algorithms
- *Recursive* Algorithms



# Iterative Algorithms

- Iterative algorithms take **one step at a time** towards the final destination (Solution).

→ Repeat operations until some condition becomes true or false.

General structure:

loop (until done)

take steps

end loop

# Iterative algorithms :Element Uniqueness

Problem:Check whether all the elements in a given array are distinct.

Unique\_Elements (A)

//Input: An array A [0..n - 1]

//Output: Returns "true" if all the elements in A are distinct and "false" otherwise.

for i  $\leftarrow$  0 to n - 2 do

for j  $\leftarrow$  i + 1 to n - 1 do

if A[i] = A [j]

return false //Not distinct

return true

# Analysis of Element Uniqueness

Worst case number of comparisons :

$$\begin{aligned}T(n) &= \sum_{i=0}^{n-2} \sum_{j=i+1}^{n-1} 1 \\&= \sum_{i=0}^{n-2} (n - 1 - (i + 1) + 1) \\&= \sum_{i=0}^{n-2} n - 1 - i \\&= \sum_{i=0}^{n-2} (n - 1) - \sum_{i=0}^{n-2} i \\&\quad \text{simplifies to:} \\&= n(n-1)/2 \\&= O(n^2)\end{aligned}$$

# Euclid's algorithm : Iterative Version

- A method for finding the **greatest common divisor** (GCD) of two integers  $m$  and  $n$ .

Example : let  $m=210$  ,  $n= 45$  .

$$\text{GCD}(210,45)=15$$

→15 is the greatest number that divides both 210 and 45.

The algorithm is based on the following fact: for  $m>n$

$$\begin{aligned}\text{GCD} ( m , n) \\ &= \text{GCD} ( n, m \bmod n) \\ &= \text{GCD} (n, m\%n)\end{aligned}$$

% is the mod operator : Returns the remainder after dividing  $m$  by  $n$

$$\begin{aligned}\text{Example : gcd}(60,24) &= \text{gcd}(24,60\%24) \\ &= \text{gcd}(24,12) \\ &= \text{gcd}(12,0) \\ &= 12\end{aligned}$$

# Euclid's Algorithm: Iterative Version

//Also known as Euclidean Algorithm.Non-recursive implementation)

```
long gcd ( long m, long n)
{
    long r;
    while (n != 0)
    {
        r = m % n;
        m = n;
        n = r;
    }
    return m;
}
```

# Tracing Euclid's Algorithm

m	n	r
210	45	30
45	30	15
30	15	0
15	0	Done .

Return m=15.

$\text{GCD}(210, 45) = 15$

# Time Complexity of Euclide -1

- How many times the loop iterates. This is the same question as:  
    How many times can we apply the mod operation?
- Each iteration reduces the first argument of the mod operation:  
    In the example we had : 210,45,30,15
- What can be the maximum number of iterations?
- Detailed analysis involves different possibilities . We are going to consider a simple approach.

# Time Complexity of Euclid -2

Consider the simple fact:  $\text{GCD}(M, N)$

If  $M > N$ , then  $(M \% N) < M/2$ . ( $r < M/2$ )

→ The size of  $M \bmod N$  is strictly less than  $M/2$ .

- So **how many times % is applied** for the remainder to be zero ( $r = 0$ )?

→ The number of times mod operations is performed  $<$  the number of times  $M$  can be divided by 2 .

→ This can not be more than  $\log_2(M)$  .

→ Worst case **Time complexity of Euclid's algorithm is  $O(\log_2 M)$** .

- **Best case** :  $O(1)$ , If at the start the remainder is 0. Ex: (100,50)

When **the worst case** occurs ?

One possibility : If  $N$  and  $M$  are two consecutive Fibonacci numbers (Check this!).



# Example : Insertion Sort

- Take multiple passes over the list
- Keep already sorted part at low-end
- Find next unsorted element
- Insert it in correct place, relative to the ones already sorted
- Invariant: each pass increases the size of sorted portion.

# Reminder : Insertion Sort Algorithm

INSERTION-SORT( $A, n$ )

**for**  $j \leftarrow 2$  **to**  $n$  **do**

$\text{key} \leftarrow A[j]$

    // Insert  $A[j]$  into the sorted sequence  $A[1 \dots j-1]$

$i \leftarrow j - 1$

**while**  $i > 0$  and  $A[i] > \text{key}$  **do**

$A[i + 1] \leftarrow A[i]$

$i \leftarrow i - 1$

$A[i + 1] \leftarrow \text{key}$

# Insertion Sort: How it Works ?

Insertion sort sorts the elements in place


input array

A: 5    2    4    6    1    3

at each iteration, the array is divided in two sub-arrays, the left subarray keeps growing

left sub-array

right sub-array



# Complexity of Insertion Sort

- The body is made up of two nested for loops. The outer loop is executed  $n - 1$  times.
- The inner loop is harder to analyze because the number of times it executes depends on how many keys in positions 1 to  $i - 1$  have a value less than the key in position  $i$ .
- In the **worst case**, each record must make its way to the start of the array. This would occur if the keys are **reversely sorted**.
- In this case, the number of comparisons will be 1, the first time through the loop, 2 the second time, and so on. Thus, the total number of comparisons will be

$$T(n) = \sum_{i=2}^n i = 2+3+\dots+(n-1) = \Theta(n^2)$$

# Complexity of Insertion Sort

- Consider the **best-case**. This occurs when the keys begin in sorted order. In this case, every pass through the inner loop will fail immediately, and **no values will be moved**.
- The total number of comparisons will be the number of times the outer loop executes:

$$\begin{aligned}T(n) &= \sum_{i=2}^n 1 = 1+1+\dots+1 \\&= (n-2+1) * 1 = n-1 \\&= \Theta(n)\end{aligned}$$

- **Average case**: We expect on average that half of the keys in the first  $i - 1$  array positions will have a value greater than that of the key at position  $i$ . Thus, the average case should be about half the cost of the worst case:

$$[n(n+1)/2] / 2 = O(n^2/4), \text{ which is still } \Theta(n^2).$$

→ Asymptotically, the average case is no better than the worst case

# Summary : Explaining Different Cases


- **The worst-case** occurs if the array is sorted in **reverse order** i.e., in decreasing order:
  - We always find that  $A[i]$  is greater than the key in the while-loop test. So, we must compare each element  $A[j]$  with each element in the entire sorted subarray  $A[1 \dots j - 1]$
- **The best case** occurs if the array is **already sorted**:
  - For each  $j = 2, 3, \dots, n$ , we find that  $A[i]$  less than or equal to the key when  $i$  has its initial value of  $(j - 1)$ . In other words, when  $i = j - 1$ , always find the key  $A[i]$  upon the first time the WHILE loop is run.
- **The average case** is also quadratic, which makes insertion sort impractical for sorting large arrays.

# Generic Insertion Sort Function: C++

```
template <class T>
void InsertionSort(T A[], int n) {
    for (int i = 1; i < n; ++i) {
        if (A[i] < A[i-1]) {
            T val = A[i];
            int j = i;
            do { A[j] = A[j-1];
                --j;
            } while ((j > 0) && (val < A[j-1]));
            A[j] = val;
        }
    }
}
```

# Binary Search(Iterative Version)

- Basic idea: On each step, look at the *middle* element of the remaining list to eliminate half of it, and quickly locate the desired element.





# Reminder : Binary Search Algorithm

*//Input A:  $a_1, a_2, \dots, a_n$ : distinct integers. Search key :  $x$*

*binary search*( $A, n, x$ )

$i \leftarrow 1$  *//left endpoint of search interval*

$j \leftarrow n$  *//right endpoint of search interval*

**while**  $i < j$  **do**

$mid \leftarrow \lfloor (i+j)/2 \rfloor$  *//midpoint*

**if**  $x > a_{mid}$  **then**

$i \leftarrow mid + 1$

**else**

$j \leftarrow mid$

**if**  $x = a_i$  **then** *//Found*

$location \leftarrow i$

**else**  $location \leftarrow 0$  *//Not found*

**return**  $location$

# Binary Search Complexity

- Best Case: match from the first comparison:  $O(1)$
- **Worst Case:** divide until reach one item, or no match
- For an array of size  $N$ , it eliminates  $\frac{1}{2}$  (half) until 1 element remains.  
 $N, N/2, N/4, N/8, \dots, 4, 2, 1$ 
  - How many divisions does it take?
- Think of it from the other direction:
  - How many times do I have to multiply by 2 to reach  $N$ ?  
 $1, 2, 4, 8, \dots, N/4, N/2, N$
  - Call this number of multiplications " $x$ ".

$$2^x = N$$

$$x = \log_2 N$$

→ Binary search is in the logarithmic complexity class.

# Worst case Analysis : Recurrence Relation

Item not in the array (size  $N$ )

$T(N)$  = number of comparisons with array elements

$$T(1) = 1$$

$$T(N) = 1 + T(N/2)$$

$$= 1 + [1 + T(N/4)]$$

$$= 2 + T(N/4)$$

$$= 2 + [1 + T(N/8)]$$

$$= 3 + T(N/8)$$

$$= \dots$$

$$= k + T(N/2^k) \quad [1]$$

# Worst-case Analysis : Recurrence

$T(N / 2^k)$  gets smaller until the base case:  $T(1)$ .

Assume

$$2^k = N$$

$$k = \log_2 N$$

Replace terms with  $k$  in [1]:

$$T(N) = \log_2 N + T(N / N)$$

$$= \log_2 N + T(1)$$

$$= \log_2 N + 1$$

→ “ $\log_2 N$ ” algorithm

# Is Binary Search more efficient?

- For a list of  $N$  elements, Binary Search can execute at most  $\log_2 N$  times.
- A **disadvantage of Binary Search** is the overhead of maintaining the list in ordered form.


N	<u>Number of Iterations</u>	
	Linear Search(Av)	Binary Search(Worst)
10	5	3
100	50	6
1,000	500	9
1,000,000	500000	20

# Binary Search: C++ Implementation

```
int binarySearch(int[] a, int target)
{
    int min = 0;
    int max = a.length - 1;
    while (min <= max) {
        int mid = (min + max) / 2;
        if (a[mid] < target)
            min = mid + 1;
        else if (a[mid] > target)
            max = mid - 1;
        else
            return mid;    // target found
    }

    return -1;    // target not found
}
```

# Recursive Algorithms



“To iterate is human, to recurse is divine!”

# Recursion vs. Iteration

- Repetition

- Iteration: explicit loop
- Recursion: repeated function calls

- Termination

- Iteration: loop condition fails
- Recursion: base case recognized
- Both can have infinite loops
- Recursive solutions can be easier to understand and to describe than iterative solutions.



## Example -1 : Recursive Euclid algorithm

```
//Base case n=0
```

```
long gcd( long m, long n)
```

```
{ if ( n == 0 )
```

```
    return m
```

```
else
```

```
    return gcd( n, m % n) }
```

Worst case complexity is the same as the iterative version :  $O(\log n)$

# Example-2: Recursive $n!$

Recursive algorithm for  $n!$

Input size:  $n$ , Basic  
operation: multiplication  
“\*”

- Analysis :Let  $M(n)$  be the **number of multiplications** needed to compute  $n!$ , then the recurrence can be expressed as:

$$M(0) = 0$$

$$M(n) = M(n-1) + 1 \quad \text{for } n > 0$$

To compute  $Factorial(n-1)$

To multiply  $Factorial(n-1)$  by  $n$

**ALGORITHM**  $Factorial(n)$

**if**  $n = 0$

**return** 1

**else**

**return**  $Factorial(n-1) * n$

# Analysis : Solving the Recurrence

$$M(0) = 0$$

$$M(n) = M(n-1) + 1 \text{ for } n > 0$$

$$\Rightarrow M(n)$$

$$= M(n-2) + 2$$

$$= M(n-3) + 3$$

$$= \dots$$

$$= M(n-n) + n$$

$$= n$$

→ The complexity of this algorithm is  $\Theta(n)$   
(Both best and worst cases)

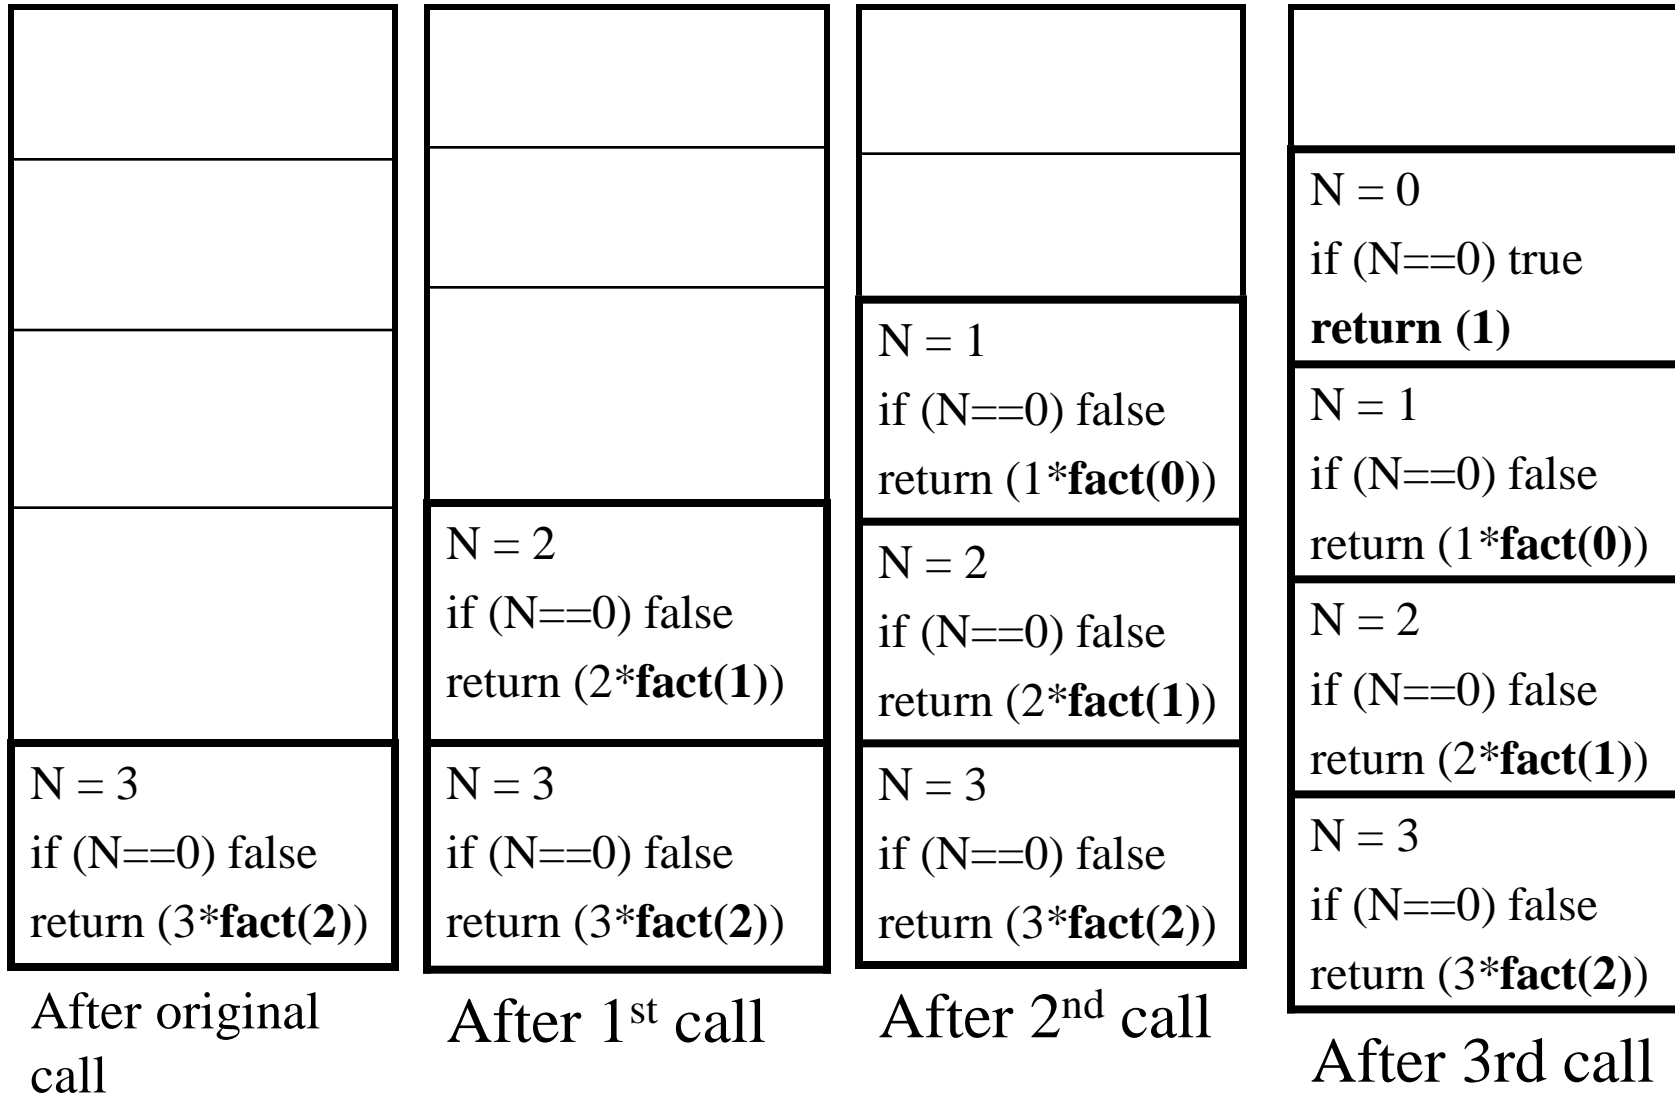
# Stack Frames in Recursion

- A **stack** is used to remember parameters and local variables across recursive calls.
- Each recursive invocation gets its own stack frame
- A stack frame contains storage for the current call :
  - Local variables
  - Parameters
  - Return info (return address and return value)
  - Other bookkeeping info
- A **new stack frame is pushed** with each recursive call
- The stack frame is popped when the function returns

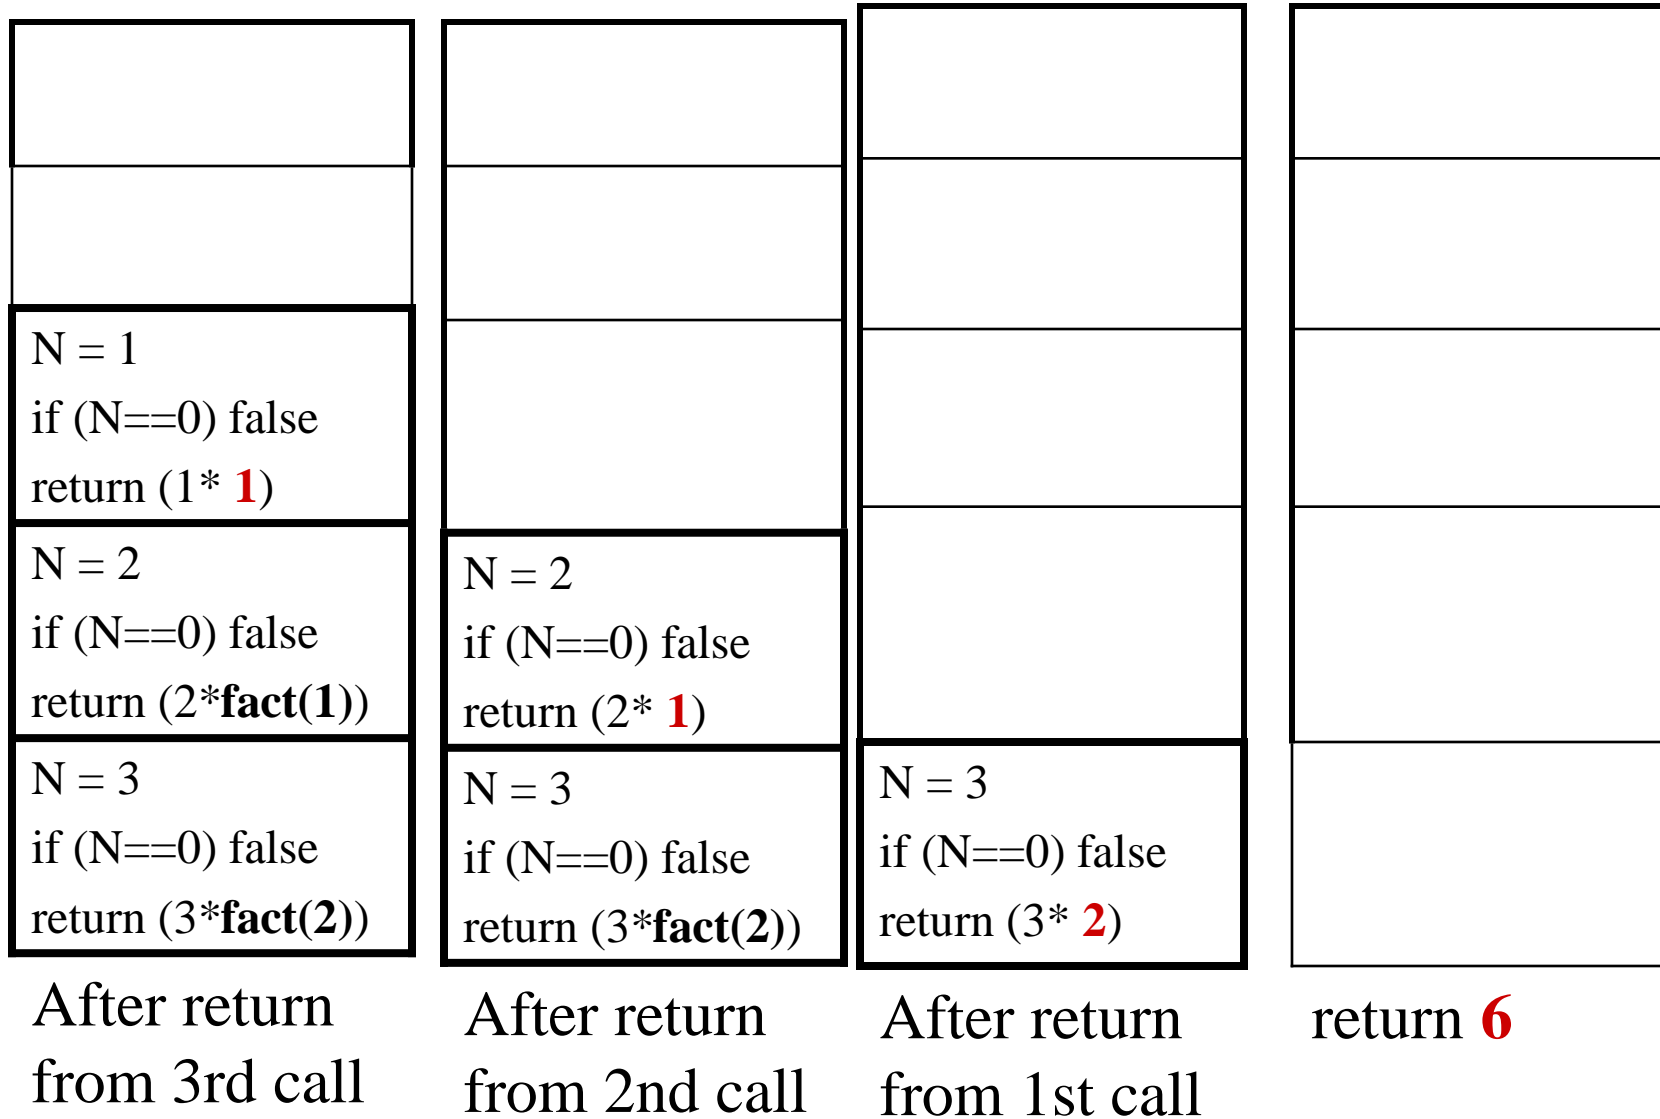
# Stack Frames in Recursion

- **When a new stack frame** is created, previous stack frames are left in the stack (In memory) .
- For example, in `fact(3)` the recursive function is called 3 times, so there will be 3 stack frames pushed on the stack.
- Each frame stores one instance and holds different values.
- As usual, **the first stack frame is at the bottom** and the last frame is at the top. The order in which the statements run is denoted by the order of stack frames.

# Tracing the call fact(3): Stack Frames



# Tracing the call fact(3): Stack Frames



# Example-3 : Recursive Fibonacci

Fibonacci sequence : 0,1, 1, 2, 3, 5, 8, 13, 21, 34, .....

// Recursively calculates nth Fibonacci number

```
long fib (int n)
```

```
{
```

```
    if( n <= 1 )
```

```
        return n
```

```
    else
```

```
        return fib(n - 1) + fib(n - 2);
```

```
}
```

This is a straightforward, but inefficient recursion ...



# Analysis of Recursive Fibonacci Algorithm

- Let  $n_k$  denote the **number of function calls** made by  $\text{Fib}(n)$ .

$$n_0 = 1 \quad \text{initial call}$$

$$n_1 = 1$$

# calls

$$n_2 = n_1 + n_0 + 1 = 1 + 1 + 1 = 3$$

$$n_3 = n_2 + n_1 + 1 = 3 + 1 + 1 = 5$$

$$n_4 = n_3 + n_2 + 1 = 5 + 3 + 1 = 9$$

$$n_5 = n_4 + n_3 + 1 = 9 + 5 + 1 = 15$$

$$n_6 = n_5 + n_4 + 1 = 15 + 9 + 1 = 25$$

$$n_7 = n_6 + n_5 + 1 = 25 + 15 + 1 = 41$$

$$n_8 = n_7 + n_6 + 1 = 41 + 25 + 1 = 67.$$


- Note that the value at least **doubles** for each two consecutive indices. Example :  **$n_4$  is more than twice  $n_2$**  ( $9 > 2 \cdot 3$ )....

$$\rightarrow n_k > 2^{k/2} \quad [\text{Ex: } n_8 > 2^{8/2}]$$

→ An exponential time algorithm!

Why again ? If not convinced , See the following recursion tree!

The recursion tree for **computing  $F_7$**  ; arrows represent recursive calls.  
Growth of the tree is very fast  
Notice repeated computations of the same value :  $F(0), F(1), F(2), F(3) \dots$



# Example-4 Recursive Binary Search

```
RecbinSearch ( A, first, last, x)
    if ( first <= last )
        mid = (first + last) / 2;    // mid is int.
        if x =A(mid )    // found it!
            return mid;
        else if x < A(mid)    // must be in 1st half
            return RecbinSearch( A, first, mid-1, x);
        else    // must be in 2nd half
            return RecbinSearch(A, mid+1, last, x);
    return -1; //Not found
```

- **No loop!** Recursive calls takes its place.

# Analysis of Recursive Binary Search

- The recurrence is the same as in Iterative algorithm  
(See the analysis slides for iterative algorithm)

$T(N)$  = Number of comparisons with array elements

$$T(1) = 1$$

$$T(N) = 1 + T(N/2)$$

.....

$$= \log_2 N + 1$$

$$= O(\log_2 N)$$

The complexity is logarithmic.

But **actual run time is bigger** in recursive binary search.

# Is Recursion Always Good ?

- Recursion works best when the algorithm and/or data structure that is used **naturally supports recursion**.
- Recursive solutions are **clearer, simpler, shorter**, and easier to understand than their non-recursive counterparts.
- From a practical software engineering point of view these are important benefits, **reducing the cost of maintaining the software**.
- However, recursive algorithms run slower than their iterative counterparts. Why?
  - When a recursive call is made, it takes time to build a **stackframe** for the call and **push it onto the system stack**. This also takes time.
  - Conversely, for every return, the **stackframe must be popped from the stack** and the local variables must be reset to their previous values — this also takes time.
- Therefore, if the recursion is deep, say, factorial(50) or Fibonacci(20), we may run out of memory!

# Multiway/ Mway Trees

What is a multiway tree ?

- A **General non-binary tree** which can be used as a search tree.
- A multiway tree differs from a binary tree in a few key ways:
  - **Each node has  $m$  children** → There are  **$m$  pointers**
  - Each node has  $m-1$  keys (Data values)
  - In nodes, keys are in ascending order :  
 $K_1 < K_2 < \dots < K_{m-1}$

# Why Multiway Trees ?


## Secondary Memory vs RAM

- We generally assume that the data would be stored in primary memory or RAM.
- Suppose the **data is too big to store in RAM**, so needs to be stored on a hard disk.
- Access time for hard disk includes;
  - *seek time + rotation time + transfer time*
- **The seek time is particularly slow** as it depends on mechanical movement : disk head physically moving to the correct position.
- The seek time is in the order of milliseconds, while CPU processes are in the order of sub-microseconds (thousands of times faster).


# Multiway /M-ary Tree:Examples

M is called the order of tree.

Example1 :  $M=5 \rightarrow$  Maximum 5 pointers,4 keys.



Example2: M=3






# Properties of Multiway Trees

Given the nonempty multiway tree, the following properties hold:

1. Each node has 0 to  $m$  subtrees(children).
2. The keys in the first  $i$  children are smaller than the  $i$ th key.
3. The keys of the data entries are **ordered in each node**.


$m = 4$ ,  $K$ :Key



# Searching an $m$ -way Search Tree

- Suppose we search an  $m$ -way search tree  $T$  for the key value  $x$ . By searching the keys of the root, we determine  $i$  such that  $K_i \leq x < K_{i+1}$ .
  - If  $x = K_i$ , the search is complete.
  - If  $x \neq K_i$ ,  $x$  must be in a subtree  $A_i$
  - We proceed to search  $x$  in subtree  $A_i$  and continue the search until we find  $x$  or determine that  $x$  is not in  $T$ .

# Search: Example-1



# B-Trees and B+ Trees

- B-Trees are balanced **m-ary** trees with two types of nodes:
  - **Leaf Nodes** and **Internal Nodes**.
  - Node capacities and total number of nodes can be very high.
- **B-trees** are **dynamic index structures** that adjust automatically to inserts and deletes.
- **B+ trees** are most widely used version of dynamic B-tree index structures. We are going to consider B+ trees.

# Properties of B+ Trees

- Can be used to form Generalized, dynamic multilevel indexes.
- Balanced Tree: The same height for paths from root to leaf.
- Interior nodes contain key values and leaf nodes contain key values and pointers to records in files.
- **Parameter:**  $n \rightarrow$  Max #of **pointers** in a node.

Every node contains:

**Key values** :  $K_i$ ,

**Pointers** :  $P_i$

# B+ tree: Parameter

- Given the parameter  $n$  (Called tree order):
  - Min # of keys in a node:  $\text{ceil}((n/2)-1)$ ,  
Max # of keys in a node:  $n-1$
  - Min # of ptrs :  $\text{ceil}(n/2)$  , Max # of ptrs:  $n$

Example :  $n=5$


#of keys in each node: 2 to 4

#of pointers “ ” : 3 to 5

- Except for the root node
  - Root node has at least one key and two pointers

# B+ Tree Index Structure

Page → Node, block



# B+Tree:Node Structures


- Data pointers stored only in **leaf nodes**.
- Internal nodes contain only keys and tree ptrs.
- Note that in any node:  $K_1 < K_2 < \dots < K_{n-1}$
- Leaf nodes form a linked(or Doubly linked)list.
- Leaf nodes contain **keys and pointers** to data file records.
- The #of levels is usually very small :  
→ Short Tree, Fast searching.



# B+ Tree :Internal Nodes


– B+ Tree is constructed by parameter **n**

- Each Node (except root) has  $\lceil n/2 \rceil$  to  $n$  pointers
- Each Node (except root) has  $\lceil n/2 \rceil - 1$  to  $n - 1$  search-key values



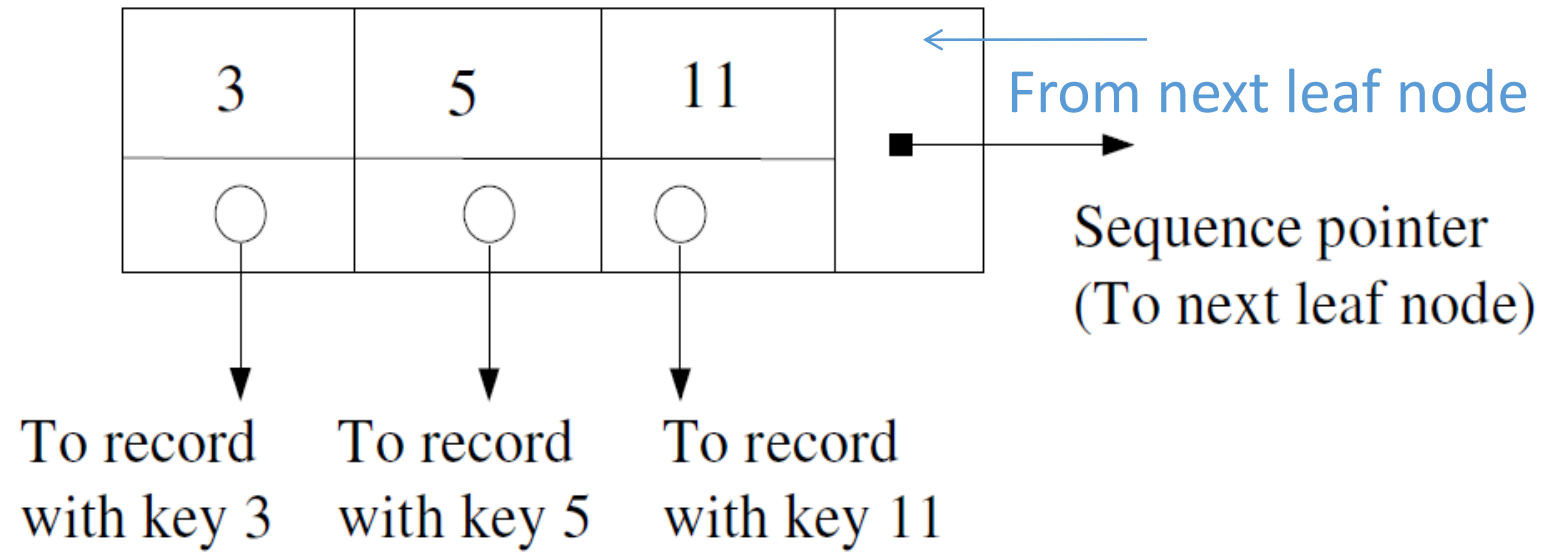
## B+ trees – internal node

---




# B+ tree – leaf node

---



# Example: B+ Tree

- 1, 4, 7, 10, 17, 21, 31, 25, 19, 20, 28, 42



# Using B+ trees: Index Search

- Suppose we have a data file and a B+ tree index to that file.
- How to access the data record whose key value is  $K_i$  ?
  - Start from the root, find the position of  $K_i$  in the root
  - Follow the pointers until a leaf node.
  - The leaf node contains  $K_i$  and ptr to the requested data record .

# Searching a B-Tree

//Search is in a node , node size ( fanout) =  $n[x]$  . A linear search is implemented

//We assume that, record pointers are stored in tree nodes.

**B-TREE-SEARCH( $x, k$ )**

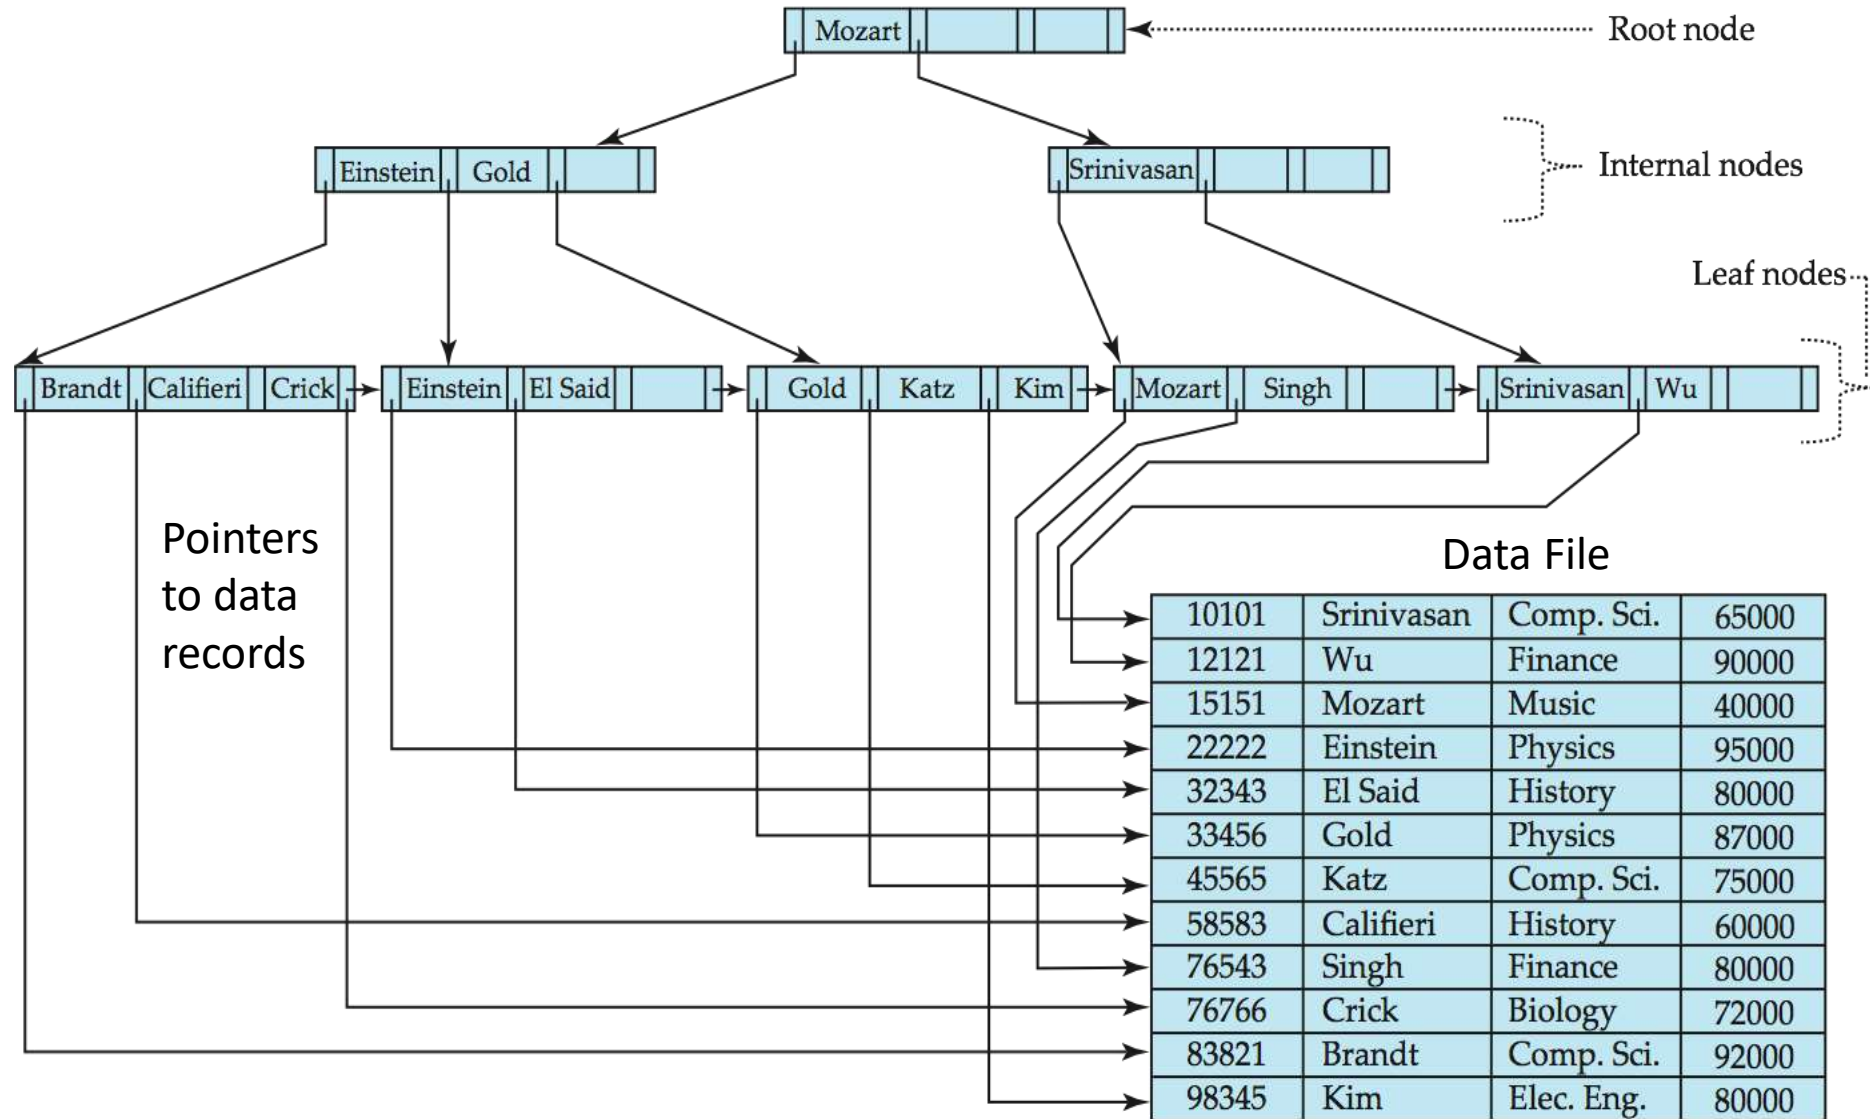
```
1   $i \leftarrow 1$ 
2  while  $i \leq n[x]$  and  $k > key_i[x]$ 
3    do  $i \leftarrow i + 1$ 
4  if  $i \leq n[x]$  and  $k = key_i[x]$ 
5    then return  $(x, i)$ 
6  if  $leaf[x]$ 
7    then return NIL
8  else DISK-READ( $c_i[x]$ )
9    return B-TREE-SEARCH( $c_i[x], k$ )
```

Start at the leftmost key in the node, and go to the right until you go too far.

If it is a leaf node, then you are done, as there is no leaf to inspect

Otherwise, retrieve the child node from the disk, and put it into memory

# An Example of B<sup>+</sup>-Tree



# Why B+ Tree Indexing ?

- B+ trees are fast,dynamic multilevel index structures.
- The root node is usually stored in main memory.
  - Locating a record using a B+ tree index requires only 3-4 disk accesses for even **very large files**!

How Efficient is the B+ indexing?

Consider an example:

Block size=4KB, Key: integer, ptr : 4B.

Node capacity= **floor**  $((4096-4)/8)$  =511 [key+ptr=8B]

Find how many index keys can be stored at **each level**.

(For node pointers,we subtract 4 or 8 from the block capacity)



# Why B+ Tree Indexing ?

- The Root : 511 keys, 512 ptrs.
- Level 1 :  $512 * 511 = 261632$  Keys  
 $512 * 512 = 262144$  ptrs
- Level 2:  $(512 * 512) * 511 = 134.000.000$  Keys !  
If this is the leaf level, we are able to index **134 million records** using just **three levels**!

How many disk accesses are needed? Three !  
level1, level2, and the data block.

→ Accessing any data record takes the same time.

Note : These are maximum possible values ,because we assume full node occupancy.

# Time Complexity of B+Tree Operations

- Remember that all leaf nodes are at the same level(Balanced).
- Assume that the parameter is n which means all nodes must be **at least half full** (n/2 keys)

- Suppose we have N data records (keys) .

1-Best case search → All Nodes are **full**.

$$T_N = \text{CEIL}(\log_{(n)} N) \quad (\text{Similar to } \log_2 N)$$

- 2-Worst case search → All nodes are only **half full**.

$$T_N = \text{CEIL}(\log_{(n/2)} N) \quad \text{Example 1: let } N=1.000.000 \text{ and } n=100$$

$$\begin{aligned} \text{Best Case : } T_{1000000} &= \text{CEIL}(\log_{100} 1.000.000) \\ &= 3 \end{aligned}$$